

# 1. Regression

## 1. Ordinary Least Squares Regression (OLSR)

**Description:** Ordinary Least Squares Regression (OLSR) is a method for estimating the unknown parameters in a linear regression model. It minimizes the sum of the squares of the differences between the observed dependent variable and those predicted by the linear function.

**Mathematics:** Given a set of data points  $(x_i, y_i)$  where  $i = 1, 2, \dots, n$ , the OLSR model is:  $y = \beta_0 + \beta_1 x + \epsilon$  where:

- $y$  is the dependent variable,
- $x$  is the independent variable,
- $\beta_0$  and  $\beta_1$  are the parameters to be estimated,
- $\epsilon$  is the error term.

The objective is to minimize:  $\sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2$

### Pros:

- Simple to understand and implement.
- Provides the best linear unbiased estimates if the assumptions are met.

### Cons:

- Assumes a linear relationship between the variables.
- Sensitive to outliers and multicollinearity.
- Assumes homoscedasticity (constant variance of errors).

### Use Cases:

- Economic forecasting.
- Risk management.
- Sales forecasting.

**How it Works:** OLSR works by fitting a line through the data points such that the sum of the squared differences between the observed values and the values predicted by the line is minimized.

## 2. Linear Regression

**Description:** Linear Regression is a statistical method to model the relationship between a dependent variable and one or more independent variables. It's the most basic form of regression.

**Mathematics:** For a single variable, the model is:  $y = \beta_0 + \beta_1 x + \epsilon$   
 For multiple variables, the model is:  $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$

The parameters  $\beta_0, \beta_1, \dots, \beta_n$  are estimated using OLS.

**Pros:**

- Easy to implement and interpret.
- Computationally efficient.

**Cons:**

- Assumes a linear relationship.
- Sensitive to outliers and multicollinearity.
- Limited to linear interactions between variables.

**Use Cases:**

- Predicting housing prices.
- Estimating product sales.
- Determining the relationship between advertising spend and revenue.

**How it Works:** Linear Regression fits a linear equation to observed data by estimating coefficients that minimize the sum of squared residuals between observed and predicted values.

### 3. Logistic Regression

**Description:** Logistic Regression is used for binary classification problems. It models the probability of a binary outcome based on one or more predictor variables.

**Mathematics:** The model is:  $\text{logit}(p) = \log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$   
 where  $p$  is the probability of the outcome being 1.

The logistic function is:  $p = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n)}}$

**Pros:**

- Suitable for binary outcomes.
- Provides probabilities and odds ratios.
- Can handle non-linear relationships through transformations.

**Cons:**

- Assumes a linear relationship between the logit of the probability and the predictor variables.
- Can be less interpretable with multiple predictors.
- Sensitive to multicollinearity.

**Use Cases:**

- Predicting disease presence (yes/no).
- Credit scoring (default/no default).
- Marketing response (purchase/no purchase).

**How it Works:** Logistic Regression estimates the probability of a binary outcome by fitting a logistic function to the data. It uses Maximum Likelihood Estimation (MLE) to find the best-fitting parameters.

## 4. Stepwise Regression

**Description:** Stepwise Regression is a method of fitting regression models in which the choice of predictive variables is carried out by an automatic procedure.

**Mathematics:** Stepwise methods include both forward selection and backward elimination:

- Forward selection starts with no variables and adds them one by one.
- Backward elimination starts with all variables and removes them one by one.

**Pros:**

- Can handle large sets of potential predictor variables.
- Simplifies models by selecting significant variables.

**Cons:**

- Can lead to overfitting.
- May miss the best model if interactions between variables are not considered.
- Sensitive to the order of variable entry/removal.

**Use Cases:**

- Model selection in epidemiology.
- Economic modeling.
- Feature selection in machine learning.

**How it Works:** Stepwise Regression iteratively adds or removes variables from the model based on specified criteria, such as the p-value or AIC/BIC, to find a subset of variables that provides the best fit.

## 5. Multivariate Adaptive Regression Splines (MARS)

**Description:** MARS is a non-parametric regression technique that can model relationships that are more complex than linear regression can handle by fitting piecewise linear regressions.

**Mathematics:** The model is:  $y = \beta_0 + \sum_{m=1}^M \beta_m h_m(x) + \epsilon$  where  $h_m(x)$  are basis functions that represent piecewise linear segments.

**Pros:**

- Can model non-linear relationships.
- Automatically selects important variables and interactions.

**Cons:**

- Can be computationally intensive.
- Requires careful tuning to avoid overfitting.
- Less interpretable than linear regression.

**Use Cases:**

- Financial forecasting.
- Environmental modeling.
- Any scenario with complex, non-linear relationships.

**How it Works:** MARS builds models by fitting piecewise linear regressions to the data, using basis functions to capture non-linearities and interactions between variables.

## 6. Locally Estimated Scatterplot Smoothing (LOESS)

**Description:** LOESS (or LOWESS) is a non-parametric method used to fit a smooth curve to a scatterplot. It is useful for visualizing the relationship between variables without assuming a specific form for the relationship.

**Mathematics:** The fitted value at point  $x$  is obtained by performing a weighted least squares regression using nearby points, where the weights decrease with distance from  $x$ .

**Pros:**

- Flexible and can model complex relationships.
- No need to specify a functional form.

**Cons:**

- Computationally intensive for large datasets.
- Sensitive to the choice of smoothing parameter.
- Difficult to interpret coefficients.

**Use Cases:**

- Trend analysis in time series data.
- Smoothing scatterplots to identify patterns.
- Exploratory data analysis.

**How it Works:** LOESS fits a smooth curve to the data by performing localized regressions. For each point, it uses nearby data points weighted by their distance to perform a weighted least squares regression, resulting in a smooth curve that captures the underlying trend.

## 2. Instance Based

### 1. k-Nearest Neighbour (kNN)

**Description:** k-Nearest Neighbour (kNN) is a non-parametric, instance-based learning algorithm used for classification and regression. It predicts the class or value of a new data point based on the majority class or average of its k-nearest neighbours in the training data.

#### Mathematics:

- **Distance Metric:** The most common distance metric used is Euclidean distance:  

$$d(x,y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$
where  $x$  and  $y$  are data points, and  $n$  is the number of features.

#### Pros:

- Simple and intuitive.
- No training phase (instance-based learning).
- Can handle multi-class classification.

#### Cons:

- Computationally expensive during prediction.
- Sensitive to irrelevant features and the choice of distance metric.
- Requires large storage for training data.

#### Use Cases:

- Pattern recognition (e.g., handwriting recognition).
- Recommender systems.
- Medical diagnosis (e.g., predicting diseases based on symptoms).

**How it Works:** kNN stores the entire training dataset. For a new data point, it calculates the distance to all training points, selects the k closest ones, and assigns the most frequent class label (classification) or the average value (regression) of these neighbors.

### 2. Learning Vector Quantization (LVQ)

**Description:** Learning Vector Quantization (LVQ) is a prototype-based supervised classification algorithm. It aims to represent the data distribution using a set of prototype vectors and classify new instances based on their proximity to these prototypes.

### Mathematics:

- **Prototype Update:** If a data point  $x$  is correctly classified by its closest prototype  $p$ , the prototype is moved closer to  $x$ . Otherwise, it is moved away.  
 $\Delta p = \alpha(x - p)$  where  $\alpha$  is the learning rate.

### Pros:

- Intuitive and interpretable prototypes.
- Can handle multi-class classification.
- Relatively simple to implement.

### Cons:

- Sensitive to the initial placement of prototypes.
- Requires careful tuning of learning rate and number of prototypes.
- May not perform well with overlapping classes.

### Use Cases:

- Image classification.
- Speech recognition.
- Medical diagnosis.

**How it Works:** LVQ initializes a set of prototype vectors. During training, each data point adjusts the closest prototype: moving it closer if correctly classified or further away if incorrectly classified. This iterative process continues until the prototypes stabilize.

## 3. Self-Organizing Map (SOM)

**Description:** Self-Organizing Map (SOM) is an unsupervised learning algorithm used for dimensionality reduction and data visualization. It maps high-dimensional data onto a lower-dimensional grid while preserving the topological relationships.

### Mathematics:

- **Weight Update:** The weight vector of the winning neuron and its neighbours are updated to move closer to the input vector.  
 $w_i(t+1) = w_i(t) + \alpha(t) h_{ci}(t) (x(t) - w_i(t))$   
where:
  - $w_i$  is the weight vector of neuron  $i$ ,
  - $\alpha(t)$  is the learning rate,
  - $h_{ci}(t)$  is the neighbourhood function.

### Pros:

- Effective for visualisation and clustering.
- Preserves topological relationships.
- Can handle noisy data.

### Cons:

- Requires careful tuning of parameters (e.g., learning rate, neighbourhood size).
- May not be suitable for large datasets.
- Training can be computationally intensive.

#### Use Cases:

- Data visualisation.
- Market segmentation.
- Image compression.

**How it Works:** SOM initialises a grid of neurons with weight vectors. For each input vector, it identifies the closest neuron (winning neuron) and updates the weights of the winning neuron and its neighbours to move closer to the input vector. This process iterates, gradually reducing the learning rate and neighbourhood size.

## 4. Locally Weighted Learning (LWL)

**Description:** Locally Weighted Learning (LWL) is a non-parametric regression method that fits a local model around the query point for each prediction. It places more weight on nearby training points.

#### Mathematics:

- **Weighting Function:** Commonly used is the Gaussian kernel:  

$$w_i = \exp\left(-\frac{d(x, x_i)^2}{2\sigma^2}\right)$$

$$w_i = \exp(-2\sigma^2 d(x, x_i)^2)$$
where  $d(x, x_i)$  is the distance between the query point  $x$  and training point  $x_i$ , and  $\sigma$  controls the width of the kernel.

#### Pros:

- Flexible and can model complex relationships.
- Naturally handles varying degrees of smoothness.
- No need for global model assumptions.

#### Cons:

- Computationally expensive, especially for large datasets.
- Sensitive to the choice of kernel and bandwidth parameter.
- Requires storing the entire training dataset.

#### Use Cases:

- Real-time prediction systems.
- Robotics (e.g., control systems).
- Financial modeling.

**How it Works:** LWL makes predictions by fitting a model to the training data weighted by their proximity to the query point. For each query, it computes weights for the training points, fits a local model using these weights, and uses the model to make the prediction.

## 5. Support Vector Machines (SVM)

**Description:** Support Vector Machines (SVM) are supervised learning algorithms used for classification and regression. They find the hyperplane that best separates the data into classes by maximizing the margin between the closest points of the classes (support vectors).

### Mathematics:

- **Optimization Problem:**  $\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2$  subject to:  $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$
- **Kernel Trick:** To handle non-linear separable data, SVM uses kernel functions (e.g., polynomial, RBF) to transform the data into higher dimensions.

### Pros:

- Effective in high-dimensional spaces.
- Can handle non-linear relationships with kernel trick.
- Robust to overfitting, especially with the right kernel.

### Cons:

- Memory-intensive and computationally expensive.
- Requires careful tuning of hyperparameters (e.g., C, kernel parameters).
- Not easily interpretable.

### Use Cases:

- Text classification (e.g., spam detection).
- Image recognition.
- Bioinformatics (e.g., protein classification).

**How it Works:** SVM works by transforming the data into a higher-dimensional space using a kernel function, if necessary, and then finding the hyperplane that maximises the margin between the classes. It uses support vectors to define this hyperplane and makes predictions based on the position of new data points relative to it.

These instance-based algorithms offer a range of techniques for classification, regression, clustering, and visualisation tasks, each with its unique strengths and limitations.

## 3.Regularisation

### 1. Ridge Regression

**Description:** Ridge Regression, also known as Tikhonov regularisation, is a technique used to analyse multiple regression data that suffer from multicollinearity. It adds a penalty



equivalent to the sum of the squared coefficients to the least squares objective function to shrink the regression coefficients.

**Mathematics:** The Ridge Regression objective function is:

$$\min_{\beta} \left( \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2 \right) \quad \text{where:}$$

- $y_i$  is the dependent variable,
- $x_{ij}$  are the independent variables,
- $\beta_j$  are the coefficients,
- $\lambda$  is the regularisation parameter.

**Pros:**

- Reduces model complexity and prevents overfitting.
- Works well with multicollinear data.
- Coefficients can be shrunk but not eliminated.

**Cons:**

- All coefficients are shrunk by the same amount, which may not be ideal.
- Selecting the regularisation parameter  $\lambda$  can be challenging.

**Use Cases:**

- Predictive modeling with highly collinear predictors.
- Financial forecasting.
- Medical research.

**How it Works:** Ridge Regression modifies the least squares objective function by adding a penalty term proportional to the square of the coefficients. This shrinks the coefficients towards zero, reducing their variance and helping to mitigate multicollinearity.

## 2. Least Absolute Shrinkage and Selection Operator (LASSO)

**Description:** LASSO is a regularisation technique that performs both variable selection and regularisation to enhance the prediction accuracy and interpretability of the statistical model it produces. It adds a penalty equivalent to the absolute value of the magnitude of coefficients to the least squares objective function.

**Mathematics:** The LASSO objective function is:

$$\min_{\beta} \left( \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j| \right)$$

**Pros:**

- Can shrink some coefficients to exactly zero, performing variable selection.
- Helps in creating simpler, more interpretable models.

- Reduces overfitting.

#### Cons:

- Can select at most  $n$  variables if there are  $n$  observations.
- The selection of the regularisation parameter  $\lambda$  is crucial.

#### Use Cases:

- Sparse models in high-dimensional data.
- Gene selection in bioinformatics.
- Economic forecasting.

**How it Works:** LASSO modifies the least squares objective function by adding a penalty term proportional to the absolute value of the coefficients. This can drive some coefficients to zero, effectively performing variable selection and reducing model complexity.

### 3. Elastic Net

**Description:** Elastic Net is a regularisation technique that linearly combines the L1 and L2 penalties of the LASSO and Ridge methods. It is particularly useful when there are multiple features which are correlated with one another.

**Mathematics:** The Elastic Net objective function is:

$$\min_{\beta} \left( \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda_1 \sum_{j=1}^p |\beta_j| + \lambda_2 \sum_{j=1}^p \beta_j^2 \right)$$

where  $\lambda_1$  and  $\lambda_2$  are the regularization parameters.

#### Pros:

- Combines the benefits of both Ridge and LASSO.
- Can handle correlated predictors well.
- Provides a more robust and flexible model.

#### Cons:

- Requires tuning of two regularisation parameters ( $\lambda_1$  and  $\lambda_2$ ).
- More complex to implement and interpret than Ridge or LASSO alone.

#### Use Cases:

- High-dimensional data with correlated features.
- Genomics and bioinformatics.
- Marketing and economics.

**How it Works:** Elastic Net modifies the least squares objective function by adding a penalty term that combines the L1 norm (LASSO) and the L2 norm (Ridge). This allows for variable selection (like LASSO) and handles multicollinearity (like Ridge).

## 4. Least-Angle Regression (LARS)

**Description:** Least-Angle Regression (LARS) is an iterative algorithm for feature selection in regression models. It is particularly efficient for high-dimensional data with many predictors.

**Mathematics:** LARS works by iteratively moving the estimated coefficients towards their least squares values, taking the largest possible step along the direction of the most correlated variable with the residuals.

### Pros:

- Efficient with high-dimensional data.
- Can produce a full piecewise linear solution path.
- Computationally less expensive than LASSO and Ridge for large datasets.

### Cons:

- Less interpretable than other regularisation methods.
- Can be sensitive to small changes in the data.

### Use Cases:

- Feature selection in large datasets.
- Genetic data analysis.
- Signal processing.

**How it Works:** LARS starts with all coefficients equal to zero and iteratively adds variables to the model. In each iteration, it identifies the variable most correlated with the residuals and takes a step in that direction. It continues this process, adjusting coefficients until all variables are included or the desired level of sparsity is achieved.

These regularisation algorithms help manage overfitting, multicollinearity, and feature selection, making them powerful tools for improving model performance and interpretability in various applications.

## 4. Decision Tree Algorithms

### 1. Classification and Regression Tree (CART)

**Description:** CART is a technique that produces either classification or regression trees, depending on whether the target variable is categorical or continuous. It is a binary recursive partitioning method, meaning it splits the data into two parts at each node.

### Mathematics:

- **Splitting Criterion:** For classification, CART uses the Gini impurity or entropy. For regression, it uses the mean squared error (MSE).

- **Gini Impurity:**  $Gini(t) = 1 - \sum_{i=1}^n p_i^2$
- **Entropy:**  $Entropy(t) = -\sum_{i=1}^n p_i \log_2(p_i)$
- **MSE:**  $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y})^2$

#### Pros:

- Easy to interpret and visualize.
- Can handle both numerical and categorical data.
- Requires little data preprocessing.

#### Cons:

- Prone to overfitting.
- Can be unstable with small changes in data.
- Greedy algorithm may not find the optimal tree.

#### Use Cases:

- Credit scoring.
- Medical diagnosis.
- Marketing analysis.

**How it Works:** CART builds the tree by recursively splitting the data into two subsets at each node, based on the criterion that minimizes impurity (for classification) or error (for regression). This process continues until a stopping condition is met (e.g., maximum depth, minimum samples per leaf).

## 2. Iterative Dichotomiser 3 (ID3)

**Description:** ID3 is an algorithm used to generate a decision tree by employing a top-down, greedy approach. It is specifically used for classification tasks.

#### Mathematics:

- **Splitting Criterion:** ID3 uses information gain, which measures the reduction in entropy.
  - **Entropy:**  $Entropy(S) = -\sum_{i=1}^n p_i \log_2(p_i)$
  - **Information Gain:**

$$Gain(S, A) = Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

$$Gain(S, A) = Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

#### Pros:

- Simple to implement.
- Results in a readable decision tree.

- Effective for small to medium-sized datasets.

#### Cons:

- Can overfit the training data.
- Only handles categorical data.
- Does not handle missing values well.

#### Use Cases:

- Customer segmentation.
- Fault diagnosis.
- Academic research.

**How it Works:** ID3 selects the attribute that maximizes information gain at each node, splitting the dataset accordingly. This process repeats recursively for each child node until all instances are classified or no further splitting is possible.

### 3. C4.5 and C5.0

**Description:** C4.5 is an extension of ID3, developed to address its limitations. C5.0 is an improved version of C4.5, offering faster performance and smaller trees.

#### Mathematics:

- **Splitting Criterion:** C4.5 uses information gain ratio, an extension of information gain that accounts for the number and size of branches.
  - **Gain Ratio:**  $\text{GainRatio}(S, A) = \frac{\text{Gain}(S, A)}{\text{SplitInfo}(A)}$
  - **SplitInfo:**  $\text{SplitInfo}(A) = -\sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \log_2 \left( \frac{|S_v|}{|S|} \right)$

#### Pros:

- Handles both categorical and continuous attributes.
- Prunes trees to avoid overfitting.
- Deals with missing values.

#### Cons:

- Can be computationally expensive.
- Prone to overfitting if not pruned properly.
- C5.0 is proprietary software.

#### Use Cases:

- Financial modeling.
- Healthcare analysis.
- Fraud detection.

**How it Works:** C4.5 builds the tree by recursively splitting the data using the attribute that maximizes the gain ratio. It then prunes the tree by removing branches that do not provide additional information. C5.0 enhances C4.5 by improving performance and reducing tree size.

## 4. Chi-squared Automatic Interaction Detection (CHAID)

**Description:** CHAID is a decision tree technique that uses chi-squared statistics to determine the best splits. It is used for both classification and regression tasks and is particularly useful for exploring relationships between categorical variables.

### Mathematics:

- **Splitting Criterion:** CHAID uses chi-squared tests to measure the independence between the target and predictor variables.
  - **Chi-squared Statistic:**  $\chi^2 = \sum \frac{(O_i - E_i)^2}{E_i}$  where  $O_i$  are observed frequencies and  $E_i$  are expected frequencies.

### Pros:

- Handles both categorical and continuous variables.
- Can produce multiway splits.
- Useful for detecting interactions between variables.

### Cons:

- Can create large and complex trees.
- Less accurate for numerical prediction.
- Sensitive to outliers.

### Use Cases:

- Market segmentation.
- Survey data analysis.
- Social science research.

**How it Works:** CHAID builds the tree by performing chi-squared tests at each node to determine the best split. It can create branches with multiple categories and merges similar categories if they do not significantly differ.

## 5. Decision Stump

**Description:** A Decision Stump is a simple decision tree with only one level, meaning it makes a decision based on a single attribute. It is often used as a weak learner in ensemble methods.

### Mathematics:

- **Splitting Criterion:** Decision Stumps can use various criteria, such as Gini impurity, entropy, or mean squared error, depending on the task (classification or regression).

#### Pros:

- Extremely simple and fast to train.
- Useful as a base model in boosting algorithms.
- Easy to interpret.

#### Cons:

- Very low predictive power.
- Can only capture very simple relationships.
- Not useful as a standalone model for complex tasks.

#### Use Cases:

- Ensemble methods (e.g., boosting).
- Baseline model for comparison.
- Simple, quick decisions in real-time systems.

**How it Works:** A Decision Stump selects the best attribute to split the data at a single level, making a decision based on that attribute alone. It is often used as a weak learner in boosting algorithms, where many stumps are combined to improve performance.

## 6. M5

**Description:** M5 is a model tree algorithm that combines decision trees and linear regression models. It builds a tree structure where each leaf node contains a linear regression model.

#### Mathematics:

- **Splitting Criterion:** M5 uses variance reduction to determine splits.
  - **Variance Reduction:**

$$\Delta = \frac{1}{N} \left[ \sum_{i=1}^N (y_i - \bar{y})^2 - \left( \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2 \right) \right]$$

$$\Delta = \frac{1}{N} \left[ \sum_{i=1}^N (y_i - \bar{y})^2 - \left( \sum_{i \in S_1} (y_i - \bar{y}_1)^2 + \sum_{i \in S_2} (y_i - \bar{y}_2)^2 \right) \right]$$

where  $\bar{y}$  is the mean of  $y$ , and  $S_1$  and  $S_2$  are the subsets after the split.

#### Pros:

- Can handle both classification and regression tasks.
- Combines the interpretability of trees with the predictive power of linear regression.
- Effective for large datasets with numerical features.

#### Cons:

- Can be complex to implement.
- Requires careful pruning to avoid overfitting.
- May be less interpretable than simpler models.

#### Use Cases:

- Financial modeling.
- Environmental modeling.
- Predictive analytics.

**How it Works:** M5 builds a tree by recursively splitting the data based on variance reduction. Each leaf node contains a linear regression model, which is used to make predictions. The tree is pruned to avoid overfitting and improve generalization.

## 7. Conditional Decision Trees

**Description:** Conditional Decision Trees (CTree) are a type of decision tree that use statistical tests to select splits, ensuring unbiased variable selection. They are part of the party package in R.

#### Mathematics:

- **Splitting Criterion:** CTree uses permutation tests to determine the association between the response and predictor variables.
  - **Test Statistic:** Based on conditional inference framework.
  - **p-values:** Computed through permutation tests to ensure unbiased variable selection.

#### Pros:

- Unbiased variable selection.
- Handles both categorical and continuous variables.
- Can model complex interactions.

#### Cons:

- Can be computationally intensive.
- Requires a sufficient sample size for accurate permutation tests.
- May be less interpretable than simpler trees.

#### Use Cases:

- Medical research.
- Social sciences.
- Market research.

**How it Works:** CTree builds the tree by performing permutation tests at each node to determine the best split. It selects the split that maximizes the association between the response and predictor variables while maintaining statistical validity.



These decision tree algorithms offer a range of techniques for classification, regression,

## 5. Bayesian Algorithms

### 1. Naive Bayes

**Description:** Naive Bayes is a probabilistic classifier based on Bayes' theorem. It assumes independence between the features given the class label, which simplifies the computation of the posterior probabilities.

#### Mathematics:

- **Bayes' Theorem:**  $P(C|x) = \frac{P(x|C) \cdot P(C)}{P(x)}$  where  $P(C|x)$  is the posterior probability of class  $C$  given features  $x$ ,  $P(x|C)$  is the likelihood,  $P(C)$  is the prior probability, and  $P(x)$  is the evidence.
- **Naive Assumption:**  $P(x|C) = \prod_{i=1}^n P(x_i|C)$

#### Pros:

- Simple and fast to train.
- Performs well with high-dimensional data.
- Robust to irrelevant features.

#### Cons:

- Assumes independence between features, which is rarely true in real-world data.
- Can be overly simplistic.
- May not perform well with small datasets.

#### Use Cases:

- Text classification (e.g., spam detection).
- Sentiment analysis.
- Document categorization.

**How it Works:** Naive Bayes calculates the posterior probability for each class using Bayes' theorem and the naive assumption of feature independence. It assigns the class with the highest posterior probability to the given instance.

### 2. Gaussian Naive Bayes

**Description:** Gaussian Naive Bayes is a variant of the Naive Bayes algorithm, specifically designed for continuous features. It assumes that the continuous features follow a Gaussian (normal) distribution.

### Mathematics:

- **Gaussian Likelihood:**  $P(x_i | C) = \frac{1}{\sqrt{2\pi}\sigma_C} \exp\left(-\frac{(x_i - \mu_C)^2}{2\sigma_C^2}\right)$  where  $\mu_C$  and  $\sigma_C$  are the mean and standard deviation of the feature  $x_i$  for class  $C$ .

### Pros:

- Handles continuous data effectively.
- Simple and computationally efficient.
- Robust to irrelevant features.

### Cons:

- Assumes a normal distribution for the features.
- Sensitive to the accuracy of the estimated mean and variance.
- Assumes independence between features.

### Use Cases:

- Real-time prediction systems.
- Medical diagnosis.
- Financial market analysis.

**How it Works:** Gaussian Naive Bayes calculates the posterior probability for each class using Bayes' theorem, assuming a Gaussian distribution for continuous features. It assigns the class with the highest posterior probability to the given instance.

## 3. Multinomial Naive Bayes

**Description:** Multinomial Naive Bayes is a variant of the Naive Bayes algorithm designed for discrete, count-based features. It is particularly suited for text classification problems.

### Mathematics:

- **Multinomial Likelihood:**  $P(x_i | C) = \frac{P(C) \cdot P(x_i | C)}{\sum_{x_i \in X} P(x_i | C)}$  where  $x_i$  represents the count of the feature, and  $X$  is the set of all features.

### Pros:

- Well-suited for text data and document classification.
- Simple and fast to train.
- Effective with high-dimensional data.

### Cons:

- Assumes feature independence.

- May not perform well with small datasets.
- Sensitive to zero counts (can be mitigated with smoothing).

#### Use Cases:

- Text classification (e.g., spam detection).
- Sentiment analysis.
- News categorization.

**How it Works:** Multinomial Naive Bayes calculates the posterior probability for each class using Bayes' theorem and the multinomial likelihood for count-based features. It assigns the class with the highest posterior probability to the given instance.

## 4. Averaged One-Dependence Estimators (AODE)

**Description:** AODE is a probabilistic classifier that extends the Naive Bayes algorithm by relaxing the independence assumption. It considers dependencies between the features and averages over several simple Bayesian classifiers.

#### Mathematics:

- **Probability Estimation:**  $P(C|x) = \frac{1}{|A|} \sum_{a \in A} P(C|x_a) \cdot P(x \setminus C, x_a) P(C \mid x) = \frac{1}{|A|} \sum_{a \in A} P(C \mid x_a) \cdot P(x \setminus C, x_a) P(C)$  where  $A$  is the set of all attributes, and  $x_a$  is the value of attribute  $a$ .

#### Pros:

- Improves accuracy by considering feature dependencies.
- Retains the simplicity and efficiency of Naive Bayes.
- Robust to irrelevant features.

#### Cons:

- Increased computational complexity compared to Naive Bayes.
- May not perform well with very large datasets.
- Can be sensitive to noisy data.

#### Use Cases:

- Text classification.
- Bioinformatics.
- Image recognition.

**How it Works:** AODE calculates the posterior probability for each class by averaging the probabilities of several simple Bayesian classifiers, each considering one feature as a parent of the others. It assigns the class with the highest averaged posterior probability to the given instance.

## 5. Bayesian Belief Network (BBN)

**Description:** Bayesian Belief Network (BBN), also known as a Bayesian Network, is a probabilistic graphical model that represents a set of variables and their conditional dependencies using a directed acyclic graph (DAG).

**Mathematics:**

- **Joint Probability Distribution:**  $P(X) = \prod_{i=1}^n P(X_i | \text{Parents}(X_i))$  where  $X$  is the set of variables, and  $\text{Parents}(X_i)$  are the parents of  $X_i$  in the network.

**Pros:**

- Captures complex dependencies between variables.
- Provides a graphical representation of relationships.
- Can handle missing data and incorporate domain knowledge.

**Cons:**

- Learning the network structure can be computationally expensive.
- Requires a large amount of data for accurate parameter estimation.
- Complex to implement and interpret.

**Use Cases:**

- Medical diagnosis.
- Risk assessment.
- Decision support systems.

**How it Works:** BBN represents the conditional dependencies between variables using a DAG. Each node in the network represents a variable, and edges represent conditional dependencies. The network is used to compute the joint probability distribution and make inferences.

## 6. Bayesian Network (BN)

**Description:** A Bayesian Network (BN) is a probabilistic graphical model that represents a set of variables and their conditional dependencies through a directed acyclic graph (DAG). It is similar to a Bayesian Belief Network but may emphasize different applications or modeling approaches.

**Mathematics:**

- **Joint Probability Distribution:**  $P(X) = \prod_{i=1}^n P(X_i | \text{Parents}(X_i))$  where  $X$  is the set of variables, and  $\text{Parents}(X_i)$  are the parents of  $X_i$  in the network.

**Pros:**

- Provides a compact representation of joint probability distributions.
- Captures and visualizes dependencies between variables.
- Can be used for both predictive and diagnostic purposes.

#### Cons:

- Structure learning can be computationally intensive.
- Requires expert knowledge for accurate modeling.
- Sensitive to the accuracy of conditional probability estimates.

#### Use Cases:

- Gene expression analysis.
- Natural language processing.
- Fraud detection.

**How it Works:** BN constructs a DAG to represent the conditional dependencies between variables. Each node represents a variable, and directed edges represent conditional dependencies. The network is used to compute the joint probability distribution and perform inference tasks.

These Bayesian algorithms offer powerful techniques for probabilistic modeling and inference, each with unique strengths and limitations tailored to different types of data and applications.

## 6. Clustering Algorithm

### 1. k-Means

**Description:** k-Means is a popular partitioning clustering algorithm that aims to divide a dataset into k clusters, where each data point belongs to the cluster with the nearest mean.

#### Mathematics:

- **Objective Function:** Minimize the sum of squared distances between each point and the centroid of its assigned cluster. 
$$\min \sum_{i=1}^k \sum_{x_j \in C_i} \|x_j - \mu_i\|^2 \quad \min \sum_{i=1}^k \sum_{x_j \in C_i} \|x_j - \mu_i\|^2$$
 where  $C_i$  is the  $i$ th cluster,  $x_j$  is a data point, and  $\mu_i$  is the centroid of cluster  $C_i$ .

#### Pros:

- Simple and fast to implement.
- Efficient for large datasets.
- Works well when clusters have a spherical shape.

#### Cons:

- Requires the number of clusters (k) to be specified in advance.

- Sensitive to the initial placement of centroids.
- Assumes clusters are of similar size and density.

#### Use Cases:

- Customer segmentation.
- Image compression.
- Document clustering.

#### How it Works:

1. Initialize k centroids randomly.
2. Assign each data point to the nearest centroid.
3. Recalculate the centroids as the mean of the points in each cluster.
4. Repeat steps 2 and 3 until convergence (no change in centroids or assignments).

## 2. k-Medians

**Description:** k-Medians is similar to k-Means but uses the median instead of the mean to update cluster centroids. This makes it more robust to outliers.

#### Mathematics:

- **Objective Function:** Minimize the sum of Manhattan distances between each point and the median of its assigned cluster. 
$$\min \sum_{i=1}^k \sum_{x_j \in C_i} \|x_j - \mu_i\|_1 \quad \min \sum_{i=1}^k \sum_{x_j \in C_i} \|x_j - \mu_i\|_1$$
 where  $\|x_j - \mu_i\|_1$  is the Manhattan distance, and  $\mu_i$  is the median of cluster  $C_i$ .

#### Pros:

- More robust to outliers than k-Means.
- Suitable for data with non-Gaussian distributions.
- Can handle clusters of different shapes.

#### Cons:

- Computationally more expensive than k-Means.
- Requires the number of clusters (k) to be specified in advance.
- Can still be sensitive to initial centroid placement.

#### Use Cases:

- Image processing.
- Market segmentation.
- Geographic clustering.

#### How it Works:

1. Initialize k medians randomly.

2. Assign each data point to the nearest median.
3. Recalculate the medians as the median of the points in each cluster.
4. Repeat steps 2 and 3 until convergence.

### 3. Expectation Maximisation (EM)

**Description:** EM is a probabilistic clustering algorithm that finds the maximum likelihood estimates of parameters in statistical models with latent variables. It is often used for Gaussian Mixture Models (GMM).

#### Mathematics:

- **Log-Likelihood:** Maximize the log-likelihood function of the observed data.  

$$\log P(X|\theta) = \sum_{i=1}^N \log \left( \sum_{j=1}^K \pi_j N(x_i | \mu_j, \Sigma_j) \right)$$

$$\log P(X|\theta) = \sum_{i=1}^N \log \left( \sum_{j=1}^K \pi_j \frac{1}{\sqrt{|\Sigma_j|}} \exp \left( -\frac{1}{2} (x_i - \mu_j)^T \Sigma_j^{-1} (x_i - \mu_j) \right) \right)$$
where  $\pi_j$  are the mixing coefficients,  $N$  is the Gaussian distribution, and  $\theta$  represents the parameters.

#### Pros:

- Can handle clusters of different shapes and sizes.
- Provides probabilistic cluster memberships.
- Can incorporate prior knowledge through Bayesian methods.

#### Cons:

- Sensitive to initial parameter estimates.
- Can converge to local maxima.
- Computationally intensive, especially for high-dimensional data.

#### Use Cases:

- Image segmentation.
- Anomaly detection.
- Natural language processing.

#### How it Works:

1. **Expectation Step (E-step):** Compute the expected value of the log-likelihood function, given the current parameter estimates.
2. **Maximization Step (M-step):** Maximize the expected log-likelihood function to update the parameter estimates.
3. Repeat the E-step and M-step until convergence.

### 4. Hierarchical Clustering

**Description:** Hierarchical clustering builds a tree (dendrogram) of clusters, where each leaf represents a data point and each node represents a cluster. It does not require specifying the number of clusters in advance.

**Mathematics:**

- **Distance Metrics:** Common metrics include Euclidean distance, Manhattan distance, and cosine similarity.
- **Linkage Criteria:** Methods to determine the distance between clusters, such as single linkage (minimum distance), complete linkage (maximum distance), and average linkage (mean distance).

**Pros:**

- No need to specify the number of clusters in advance.
- Produces a hierarchy of clusters.
- Suitable for small to medium-sized datasets.

**Cons:**

- Computationally expensive for large datasets.
- Sensitive to noise and outliers.
- Choice of distance metric and linkage criterion can affect results.

**Use Cases:**

- Gene expression analysis.
- Document clustering.
- Social network analysis.

**How it Works:**

1. **Agglomerative (Bottom-Up) Approach:**
  - Start with each data point as a separate cluster.
  - Iteratively merge the closest clusters based on a chosen distance metric and linkage criterion.
  - Continue until all points are in a single cluster.
2. **Divisive (Top-Down) Approach:**
  - Start with all data points in a single cluster.
  - Iteratively split the cluster into smaller clusters.
  - Continue until each point is its own cluster.

These clustering algorithms offer diverse approaches to partitioning data, each with unique strengths and limitations suited to different types of data and applications.



# 7. Association Rule Learning Algorithms

## 1. Apriori Algorithm

**Description:** The Apriori algorithm is used for discovering frequent itemsets in transactional databases and generating association rules between items based on their frequency of co-occurrence.

### Mathematics:

- **Support:** Measures the frequency of an itemset in the dataset.  
$$\text{Support}(A \rightarrow B) = \frac{\text{transactions containing both } A \text{ and } B}{\text{total transactions}}$$
- **Confidence:** Measures the likelihood of item B being purchased when item A is purchased.  
$$\text{Confidence}(A \rightarrow B) = \frac{\text{Support}(A \rightarrow B)}{\text{Support}(A)}$$
- **Apriori Property:** If an itemset is frequent, then all of its subsets must also be frequent.

### Pros:

- Simple and easy to understand.
- Effective for large datasets.
- Provides quantitative measures (support, confidence).

### Cons:

- Computationally expensive, especially for large itemsets.
- Generates a large number of rules, including many irrelevant ones.
- Requires tuning of minimum support and confidence thresholds.

### Use Cases:

- Market basket analysis.
- Recommender systems.
- Cross-selling and upselling strategies.

### How it Works:

1. **Generate Candidate Itemsets:**
  - Start with frequent individual items (1-itemsets).
  - Generate candidate itemsets of higher sizes based on frequent (k-1)-itemsets.
2. **Calculate Support:**
  - Count the occurrences of each candidate itemset in the dataset.
  - Discard itemsets that do not meet the minimum support threshold.

### 3. **Generate Association Rules:**

- Generate rules that meet the minimum confidence threshold from the frequent itemsets.

## 2. **Eclat Algorithm**

**Description:** The Eclat (Equivalence Class Transformation) algorithm is another method for discovering frequent itemsets in transactional databases. It differs from Apriori in that it uses a vertical data format (transaction ID/itemset) instead of a horizontal format (item/transaction ID).

### **Mathematics:**

- **Transaction ID/itemset structure:** Lists items against transaction IDs.
- **Support:** Defined similarly to Apriori, measuring the frequency of an itemset in the dataset.

### **Pros:**

- Efficient in terms of memory and computation.
- Particularly effective for datasets with a large number of transactions but smaller itemsets.
- No need to generate candidate itemsets of higher sizes.

### **Cons:**

- Cannot handle datasets with large itemsets as efficiently as Apriori.
- Less intuitive in terms of rule generation compared to Apriori.
- May not be suitable for datasets with sparse transactions.

### **Use Cases:**

- Market basket analysis in large retail datasets.
- Web usage mining.
- Analysis of clickstream data.

### **How it Works:**

1. **Vertical Data Format:**
  - Organize the dataset into a vertical format (transaction ID/itemset).
2. **Generate Equivalence Classes:**
  - Identify all itemsets that share at least one item.
3. **Calculate Support:**
  - Count the number of transactions in each equivalence class.
4. **Generate Frequent Itemsets:**
  - Determine frequent itemsets that meet the minimum support threshold.

Both Apriori and Eclat algorithms are essential tools for discovering associations and patterns in large datasets. The choice between them often depends on the specific characteristics of the dataset, such as size, sparsity, and the distribution of itemsets.

## 8. Dimensionality Reduction Algorithms

### 1. Principal Component Analysis (PCA)

**Description:** PCA is a linear dimensionality reduction technique that transforms the data into a new coordinate system such that the greatest variance by any projection of the data comes to lie on the first coordinate, the second greatest variance on the second coordinate, and so on.

#### Mathematics:

- **Eigenvalue Decomposition:** PCA finds the eigenvectors and eigenvalues of the covariance matrix of the data.
- **Variance Maximization:** It maximizes the variance of projected data along the principal components.

#### Pros:

- Reduces the dimensionality of data while preserving as much variance as possible.
- Speeds up learning algorithms.
- Removes correlated features.

#### Cons:

- Assumes linear relationships between variables.
- May not perform well on non-linear data.
- Interpreting principal components can be complex.

#### Use Cases:

- Image compression.
- Exploratory data analysis.
- Feature extraction in machine learning.

#### How it Works:

1. **Compute Covariance Matrix:** Calculate the covariance matrix of the dataset.
2. **Eigenvalue Decomposition:** Compute the eigenvectors and eigenvalues of the covariance matrix.
3. **Select Principal Components:** Choose the top  $k$  eigenvectors corresponding to the largest eigenvalues to form the new feature subspace.

### 2. Principal Component Regression (PCR)

**Description:** PCR is a regression technique that uses PCA to reduce the number of predictor variables in a regression model to address issues of multicollinearity.

#### Mathematics:

- **Regression with Principal Components:** Use PCA to transform predictors, then perform regression on the transformed data.

#### Pros:

- Mitigates multicollinearity in regression.
- Simplifies the model by reducing the number of predictors.

#### Cons:

- Loss of interpretability of original predictors.
- Requires careful interpretation of principal components.

#### Use Cases:

- Regression with highly correlated predictors.
- Predictive modeling with high-dimensional data.

#### How it Works:

1. **Apply PCA:** Reduce predictor variables using PCA.
2. **Perform Regression:** Fit a regression model using the reduced set of principal components as predictors.

### 3. Partial Least Squares Regression (PLSR)

**Description:** PLSR is a technique that extends regression by finding a linear regression model by projecting the predicted variables and the observable variables to a new space.

#### Mathematics:

- **Maximize Covariance:** PLSR iteratively extracts components that maximize covariance between predictors and response variables.

#### Pros:

- Handles multicollinearity and small sample sizes well.
- Incorporates response variable information into the dimensionality reduction process.

#### Cons:

- Can overfit small datasets.
- Less interpretable compared to standard regression models.

#### Use Cases:

- Predictive modeling in chemometrics.
- Marketing research.
- Bioinformatics.

#### How it Works:

1. **Iterative Component Extraction:** Sequentially extract components that maximize covariance between predictors and response variables.
2. **Model Fitting:** Fit a regression model using the extracted components.

## 4. Sammon Mapping

**Description:** Sammon Mapping is a non-linear dimensionality reduction technique that aims to preserve the local structure of the data in the reduced-dimensional space.

#### Mathematics:

- **Stress Function:** Minimizes the stress function, which measures the mismatch between pairwise distances in the original and reduced spaces.

#### Pros:

- Preserves local relationships in the data.
- Suitable for non-linear relationships.

#### Cons:

- Computationally intensive.
- Sensitive to noise and outliers.

#### Use Cases:

- Pattern recognition.
- Clustering analysis.
- Visualization of high-dimensional data.

#### How it Works:

1. **Calculate Dissimilarities:** Compute dissimilarities (distances) between data points.
2. **Optimize Stress Function:** Minimize the stress function using iterative optimization techniques.
3. **Embedding:** Map high-dimensional data into a lower-dimensional space while preserving local structure.

## 5. Multidimensional Scaling (MDS)

**Description:** MDS is a technique that constructs a low-dimensional representation of data while preserving the distances (or dissimilarities) between data points as much as possible.

#### Mathematics:

- **Stress Function:** Minimizes the discrepancy between pairwise distances in the original and reduced spaces.

**Pros:**

- Provides a visual representation of similarity relationships.
- Preserves pairwise distances.

**Cons:**

- Computational complexity increases with dataset size.
- Sensitive to noise and outliers.

**Use Cases:**

- Social sciences (e.g., psychology, sociology).
- Marketing research.
- Bioinformatics.

**How it Works:**

1. **Distance Matrix:** Construct a distance matrix from the original data.
2. **Optimize Stress Function:** Minimize the stress function using iterative methods (e.g., gradient descent).
3. **Embedding:** Project data points into a lower-dimensional space while preserving pairwise distances.

## 6. Projection Pursuit

**Description:** Projection Pursuit is a technique that seeks to find interesting projections of high-dimensional data by optimizing a criterion that measures the "interestingness" of projections.

**Mathematics:**

- **Projection Index:** Maximizes or minimizes a projection index that quantifies the interestingness of projections.

**Pros:**

- Uncovers interesting and informative projections.
- Can handle non-linear relationships.

**Cons:**

- Subjective choice of projection index.
- Computationally intensive.

**Use Cases:**

- Anomaly detection.
- Exploratory data analysis.
- Visualization of complex datasets.

### How it Works:

1. **Define Projection Index:** Choose a projection index that measures the "interestingness" of projections.
2. **Optimize Projection Index:** Maximize or minimize the chosen index using optimization techniques.
3. **Projection:** Project data into a lower-dimensional space based on the optimized index.

## 7. Linear Discriminant Analysis (LDA)

**Description:** LDA is a supervised dimensionality reduction technique that finds linear combinations of features that best separate two or more classes in the data.

### Mathematics:

- **Maximize Between-Class Scatter:** Maximizes the ratio of between-class scatter to within-class scatter.

### Pros:

- Maximizes class separability.
- Provides insights into class differences.

### Cons:

- Assumes linear separability.
- Requires labeled data.

### Use Cases:

- Face recognition.
- Speech recognition.
- Biometric identification.

### How it Works:

1. **Compute Class Means and Scatter:** Calculate mean vectors and scatter matrices for each class.
2. **Optimize Criterion:** Maximize the ratio of between-class scatter to within-class scatter.
3. **Project Data:** Project data into a lower-dimensional space using the derived linear discriminants.

## 8. Mixture Discriminant Analysis (MDA)

**Description:** MDA is an extension of LDA that generalizes to multiple classes and considers mixtures of Gaussian distributions for each class.

### Mathematics:

- **Gaussian Mixture Models:** Represents each class as a mixture of Gaussian distributions.

**Pros:**

- Handles complex class distributions.
- Effective for multi-class problems.

**Cons:**

- Computationally intensive.
- Requires estimation of Gaussian parameters.

**Use Cases:**

- Pattern recognition.
- Medical diagnostics.
- Financial forecasting.

**How it Works:**

1. **Fit Gaussian Mixtures:** Estimate parameters of Gaussian mixtures for each class.
2. **Compute Class Conditional Densities:** Calculate class conditional densities using the Gaussian mixtures.
3. **Optimize Discriminant Functions:** Maximize likelihood or posterior probability to derive discriminant functions.

## 9. Quadratic Discriminant Analysis (QDA)

**Description:** QDA is a non-linear extension of LDA that assumes each class has its own covariance matrix, allowing for more complex decision boundaries.

**Mathematics:**

- **Class-Specific Covariance Matrices:** Each class has its own covariance matrix.

**Pros:**

- More flexible decision boundaries than LDA.
- Handles non-linear relationships.

**Cons:**

- Requires more data than LDA.
- Sensitive to overfitting.

**Use Cases:**

- Image classification.
- Natural language processing.



- Financial risk modeling.

#### How it Works:

1. **Estimate Covariance Matrices:** Compute covariance matrices for each class.
2. **Compute Discriminant Functions:** Derive quadratic discriminant functions based on class-specific covariance matrices.
3. **Classify Data:** Assign class labels based on the discriminant functions.

## 10. Flexible Discriminant Analysis (FDA)

**Description:** FDA is a flexible extension of linear and quadratic discriminant analysis that allows for the incorporation of non-linear transformations of predictors.

#### Mathematics:

- **Non-linear Transformations:** Applies non-linear transformations to predictors.

#### Pros:

- Captures complex relationships between predictors and response variables.
- Improves classification accuracy over linear methods.

#### Cons:

- Requires more computational resources.
- May lead to overfitting with small datasets.

#### Use Cases:

- Complex pattern recognition.
- Functional data analysis.
- Bioinformatics.

#### How it Works:

1. **Non-linear Transformation:** Apply non-linear transformations to predictors.
2. **Derive Discriminant Functions:** Construct discriminant functions using transformed predictors.
3. **Classify Data:** Assign class labels based on discriminant scores.

# Deep learning

## 1. Convolutional Neural Networks (CNNs)

**Description:** CNNs are specialized neural networks designed to process structured grid data, such as images. They use convolutional layers to capture spatial hierarchies in the data.

**Mathematics:**

- **Convolution Operation:**  $(I * K)(i, j) = \sum_m \sum_n I(i-m, j-n) K(m, n)$  where  $I$  is the input image,  $K$  is the kernel, and  $*$  denotes the convolution operation.

**Use Cases:**

- Image classification (e.g., ImageNet)
- Object detection (e.g., YOLO, Faster R-CNN)
- Image segmentation (e.g., U-Net)

## 2. Recurrent Neural Networks (RNNs)

**Description:** RNNs are neural networks designed to recognize patterns in sequences of data. They maintain a hidden state that captures information about previous time steps.

**Mathematics:**

- **Hidden State Update:**  $h_t = f(W_h h_{t-1} + W_x x_t + b)$  where  $h_t$  is the hidden state at time  $t$ ,  $x_t$  is the input at time  $t$ , and  $f$  is an activation function.

**Use Cases:**

- Language modeling
- Speech recognition
- Time series forecasting

## 3. Long Short-Term Memory Networks (LSTMs)

**Description:** LSTMs are a type of RNN designed to better capture long-term dependencies in sequences. They use gates to control the flow of information.

**Mathematics:**

- **LSTM Cell:** 
$$\begin{aligned} f_t &= \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \\ i_t &= \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \\ C_t &= \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \\ o_t &= \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \end{aligned}$$
$$\tilde{C}_t = f_t * C_{t-1} + i_t * C_t$$
$$C_t = \tilde{C}_t * C_{t-1} + i_t * C_t$$
$$h_t = f_t * C_{t-1} + i_t * C_t$$

$$\begin{aligned}
& + i_t * \tilde{C}_t \\
C_t &= f_t * C_{t-1} + i_t * \tilde{C}_t \\
o_t &= \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \\
h_t &= o_t * \tanh(C_t) \\
h_t &= o_t * \tanh(C_t)
\end{aligned}$$

#### Use Cases:

- Text generation
- Machine translation
- Handwriting recognition

## 4. Generative Adversarial Networks (GANs)

**Description:** GANs consist of two neural networks, a generator and a discriminator, which are trained adversarially. The generator creates fake data, while the discriminator evaluates its authenticity.

#### Mathematics:

- **Generator Loss:**  $\mathcal{L}_G = -\log(D(G(z)))$
- **Discriminator Loss:**  $\mathcal{L}_D = -\left[\log(D(x)) + \log(1 - D(G(z)))\right]$

#### Use Cases:

- Image generation
- Text-to-image synthesis
- Data augmentation

## 5. Transformer Networks

**Description:** Transformers are models that use self-attention mechanisms to process sequences. They excel in natural language processing tasks due to their ability to handle long-range dependencies.

#### Mathematics:

- **Self-Attention:**  $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$   
where  $Q$  is the query,  $K$  is the key,  $V$  is the value, and  $d_k$  is the dimension of the key.

#### Use Cases:

- Machine translation
- Text summarization
- Language modeling

## 6. Autoencoders

**Description:** Autoencoders are neural networks designed to learn efficient representations of data, typically for the purpose of dimensionality reduction or denoising.

**Mathematics:**

- **Encoding:**  $z = f(Wx + b)$
- **Decoding:**  $\hat{x} = g(W'z + b')$  where  $z$  is the latent representation,  $\hat{x}$  is the reconstruction, and  $f$  and  $g$  are activation functions.

**Use Cases:**

- Anomaly detection
- Image denoising
- Dimensionality reduction

## 7. Deep Belief Networks (DBNs)

**Description:** DBNs are generative models composed of multiple layers of Restricted Boltzmann Machines (RBMs). They can learn hierarchical representations of data.

**Mathematics (RBM):**

- **Energy Function:**  $E(v, h) = -\sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} v_i w_{ij} h_j$  where  $v$  is the visible layer,  $h$  is the hidden layer,  $a$  and  $b$  are biases, and  $w$  are weights.

**Use Cases:**

- Feature learning
- Dimensionality reduction
- Pre-training for deep networks

## 8. Deep Q-Networks (DQNs)

**Description:** DQNs are a type of deep reinforcement learning algorithm where a neural network approximates the Q-value function, enabling agents to learn policies for decision-making tasks.

**Mathematics:**

- **Q-Value Update:**  $Q(s, a) = Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$  where  $s$  is the state,  $a$  is the action,  $r$  is the reward,  $\gamma$  is the discount factor, and  $\alpha$  is the learning rate.

**Use Cases:**

- Game playing (e.g., Atari games)

- Robotics
- Autonomous driving

## 9. Variational Autoencoders (VAEs)

**Description:** VAEs are generative models that learn to encode data into a latent space and decode from this space to reconstruct data. They introduce a probabilistic approach to autoencoders.

**Mathematics:**

- **Loss Function:**  $L = \mathbb{E}_{q(z|x)} [\log p(x|z)] - \text{DKL}(q(z|x) \| p(z))$  where  $\text{DKL}$  is the Kullback-Leibler divergence.

**Use Cases:**

- Image generation
- Data imputation
- Anomaly detection

## 10. Graph Neural Networks (GNNs)

**Description:** GNNs are neural networks that operate on graph-structured data, learning to represent nodes, edges, or entire graphs.

**Mathematics:**

- **Message Passing:**  $h_v^{(k)} = \sigma \left( \sum_{u \in N(v)} W h_u^{(k-1)} + b \right)$  where  $h_v^{(k)}$  is the hidden state of node  $v$  at layer  $k$ ,  $N(v)$  are the neighbors of  $v$ , and  $W$  and  $b$  are learnable parameters.

**Use Cases:**

- Social network analysis
- Molecular structure analysis
- Recommender systems

## 1. Radial Basis Function Networks (RBFNs)

**Description:** RBFNs are a type of artificial neural network that uses radial basis functions as activation functions. They are typically used for function approximation, classification, and regression tasks.

**Mathematics:**

- **Radial Basis Function:**  $\phi(r) = \phi(\|x - c\|)$  where  $x$  is the input,  $c$  is the center of the radial basis function, and  $\|x - c\|$  is the Euclidean distance.
- **Network Output:**  $y(x) = \sum_{i=1}^N w_i \phi(\|x - c_i\|)$  where  $w_i$  are the weights,  $c_i$  are the centers, and  $N$  is the number of neurons in the hidden layer.

#### Use Cases:

- Time series prediction
- Function approximation
- Pattern recognition

## 2. Multilayer Perceptrons (MLPs)

**Description:** MLPs are feedforward neural networks consisting of at least three layers: an input layer, one or more hidden layers, and an output layer. They use non-linear activation functions to model complex patterns in data.

#### Mathematics:

- **Forward Pass:**  $a(l) = f(W(l)a(l-1) + b(l))$  where  $W$  is the weight matrix,  $b$  is the bias vector,  $a$  is the activation vector, and  $f$  is the activation function.
- **Backpropagation:**  $\delta(l) = \partial L / \partial z(l)$  where  $\delta(l) = \partial z(l) / \partial L$  and  $\Delta W(l) = \delta(l)(a(l-1))^T$  where  $\delta$  is the error term, and  $L$  is the loss function.

#### Use Cases:

- Image and speech recognition
- Predictive modeling
- Classification and regression tasks

## 3. Self-Organizing Maps (SOMs)

**Description:** SOMs are unsupervised learning algorithms that produce a low-dimensional, discretized representation of the input space, called a map. They are used for clustering and visualization.

#### Mathematics:

- **Distance Calculation:** Compute the distance between input vector  $x$  and weight vector  $w$ .  $d = \|x - w\|$
- **Update Rule:**  $w(t+1) = w(t) + \eta(t) \cdot h(i, t) \cdot (x - w(t))$  where  $\eta$  is the learning rate, and  $h$  is the neighborhood function.

#### Use Cases:

- Customer segmentation
- Pattern recognition
- Dimensionality reduction

## 4. Restricted Boltzmann Machines (RBMs)

**Description:** RBMs are generative stochastic neural networks that can learn a probability distribution over a set of inputs. They consist of visible and hidden units with a symmetric connection between them.

**Mathematics:**

- **Energy Function:**  $E(v, h) = -\sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} v_i w_{ij} h_j$   
 $E(v, h) = -\sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} v_i w_{ij} h_j$   
 where  $v$  is the visible layer,  $h$  is the hidden layer,  $a$  and  $b$  are biases, and  $w$  are weights.

**Use Cases:**

- Feature learning
- Dimensionality reduction
- Collaborative filtering

## 5. Recursive Neural Networks

**Description:** Recursive Neural Networks (RvNNs) are a type of deep neural network that applies the same set of weights recursively over a structured input to produce a structured prediction. They are commonly used for parsing hierarchical structures such as sentences or images.

**Mathematics:**

- **Recursive Composition:**  $h = f(W[h_1, h_2, \dots, h_k] + b)$   
 $h = f(W[h_1, h_2, \dots, h_k] + b)$  where  $h$  is the hidden state,  $W$  is the weight matrix,  $b$  is the bias vector, and  $f$  is a non-linear activation function.

**Use Cases:**

- Natural language processing (e.g., parsing, sentiment analysis)
- Image segmentation
- Scene parsing

## 6. Backpropagation

**Description:** Backpropagation is a supervised learning algorithm used for training artificial neural networks. It calculates the gradient of the loss function with respect to each weight by the chain rule, updating the weights to minimize the loss.

**Mathematics:**

- **Loss Function Gradient:**  $\delta_j(L) = \frac{\partial L}{\partial z_j(L)} \delta_j(L) = \frac{\partial L}{\partial z_j(L)} \delta_j(L)$   
 $\delta_j(L) = \frac{\partial L}{\partial z_j(L)} \delta_j(L) = \frac{\partial L}{\partial z_j(L)} \delta_j(L)$   
 $(a^{(l)})^T \Delta W(l) = \delta^{(l+1)} (a^{(l)})^T$  where  $\delta_j(L)$  is the error term for the output layer,  $\Delta W(l)$  is the weight update for layer  $l$ , and  $L$  is the loss function.

#### Use Cases:

- Training neural networks (e.g., MLPs, CNNs, RNNs)
- Optimizing model parameters
- Any supervised learning task with neural networks

## Feedforward Neural Networks (FNNs)

**Description:** Feedforward Neural Networks (FNNs) are the simplest type of artificial neural network, where connections between the nodes do not form a cycle. They consist of an input layer, one or more hidden layers, and an output layer. Information moves in one direction—from input to output.

**Forward Propagation:** Forward propagation is the process by which input data passes through the network, layer by layer, until it reaches the output layer. During this process, each layer performs a set of transformations on the input data.

#### Mathematics:

- **Input Layer:** Let  $x$  be the input vector.
- **Hidden Layer (i-th layer):**  $a^{(i)} = f(W^{(i)}z^{(i-1)} + b^{(i)})$  where:
  - $z^{(i-1)}$  is the output from the  $(i-1)$ -th layer (for the input layer,  $z^{(0)} = x$ ),
  - $W^{(i)}$  is the weight matrix for the  $i$ -th layer,
  - $b^{(i)}$  is the bias vector for the  $i$ -th layer,
  - $f$  is the activation function (e.g., ReLU, Sigmoid, Tanh),
  - $a^{(i)}$  is the activated output of the  $i$ -th layer.
- **Output Layer:** The output layer transforms the final hidden layer's activations into the network's output.  $y = g(W^{(L)}z^{(L-1)} + b^{(L)})$  where:
  - $y$  is the output vector,
  - $g$  is the activation function of the output layer (often softmax for classification, linear for regression).

#### Use Cases:

- Classification (e.g., image and text classification)
- Regression (e.g., predicting continuous values)
- Function approximation

## Forward Propagation Process



1. **Input Data:** The input data  $x$  is fed into the input layer.
2. **Hidden Layers:** Each hidden layer processes the input data from the previous layer:
  - Calculate the weighted sum of the inputs and biases.
  - Apply the activation function to introduce non-linearity.
3. **Output Layer:** The final hidden layer's output is processed by the output layer to generate the network's prediction.

**Example:** Suppose an FNN with one hidden layer:

- Input layer:  $x$
- Hidden layer:  $h$
- Output layer:  $y$

Let:

- $W_1$  and  $b_1$  be the weights and biases for the hidden layer.
- $W_2$  and  $b_2$  be the weights and biases for the output layer.

The forward propagation steps are:

1. Hidden layer:  $h = f(W_1x + b_1)$
2. Output layer:  $y = g(W_2h + b_2)$

where  $f$  and  $g$  are activation functions.

## Pros and Cons

**Pros:**

- **Simplicity:** Easy to understand and implement.
- **Flexibility:** Can be used for a variety of tasks, including classification and regression.
- **Efficiency:** Suitable for problems where data is not sequential or structured.

**Cons:**

- **Overfitting:** Prone to overfitting, especially with small datasets or too many hidden layers.
- **Scalability:** May not perform well with very large and complex datasets.
- **No Memory:** Cannot handle sequential data as it does not have any memory of past inputs (unlike RNNs or LSTMs).

**Use Cases:**

- Image recognition
- Text classification
- Predictive analytics
- Medical diagnosis (e.g., disease prediction)
- Financial forecasting

